


Modern Operating Systems

Fifth Edition



Chapter 2


Threads

Threads


- Implicit assumption so far:
 - 1 process → 1 thread of execution
- Multithreaded execution:
 - 1 process → N threads of execution
- Why allow multiple **threads** per process?
 - Lightweight processes
 - Allow space- and time- efficient parallelism
 - Organized in thread groups
 - Allow simple communication and synchronization

Concurrency vs. Parallelism

□ Concurrent execution on single-core system:



□ Parallelism on a multi-core system:




Threads Usage

- ❑ Many web browsers ran as single process (some still do)
 - ❑ If one web site causes trouble, entire browser can hang or crash
- ❑ Google Chrome Browser is multiprocess with 3 different types of processes:
 - ❑ **Browser** process manages user interface, disk and network I/O
 - ❑ **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - ❑ **Plug-in** process for each type of plug-in



Thread Usage

A word processor with three threads.




Thread Benefits

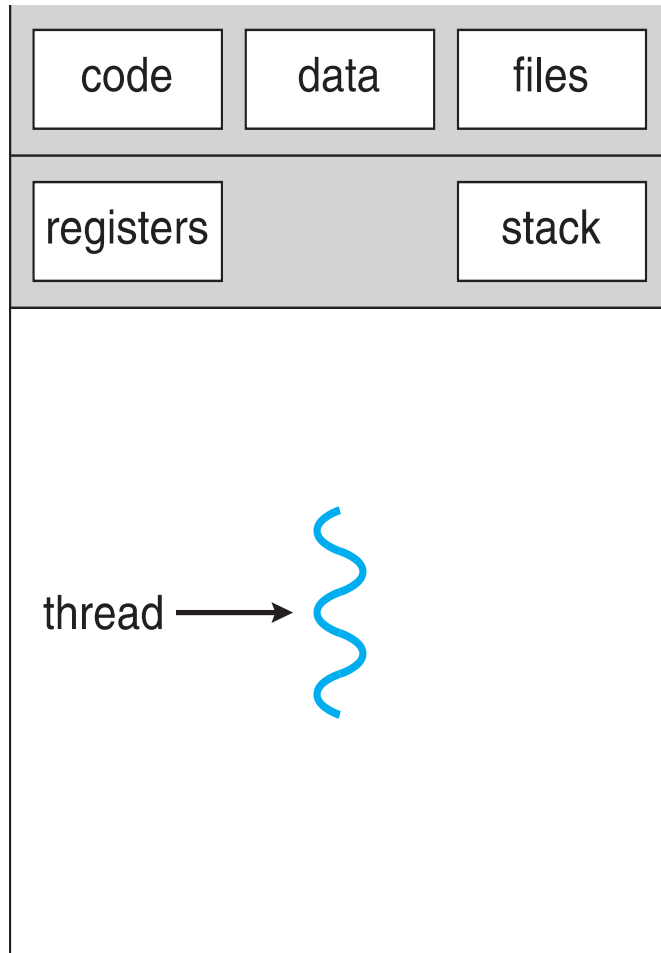
- ❑ **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- ❑ **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- ❑ **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- ❑ **Scalability** – process can take advantage of multiprocessor architectures

Threads and Processes

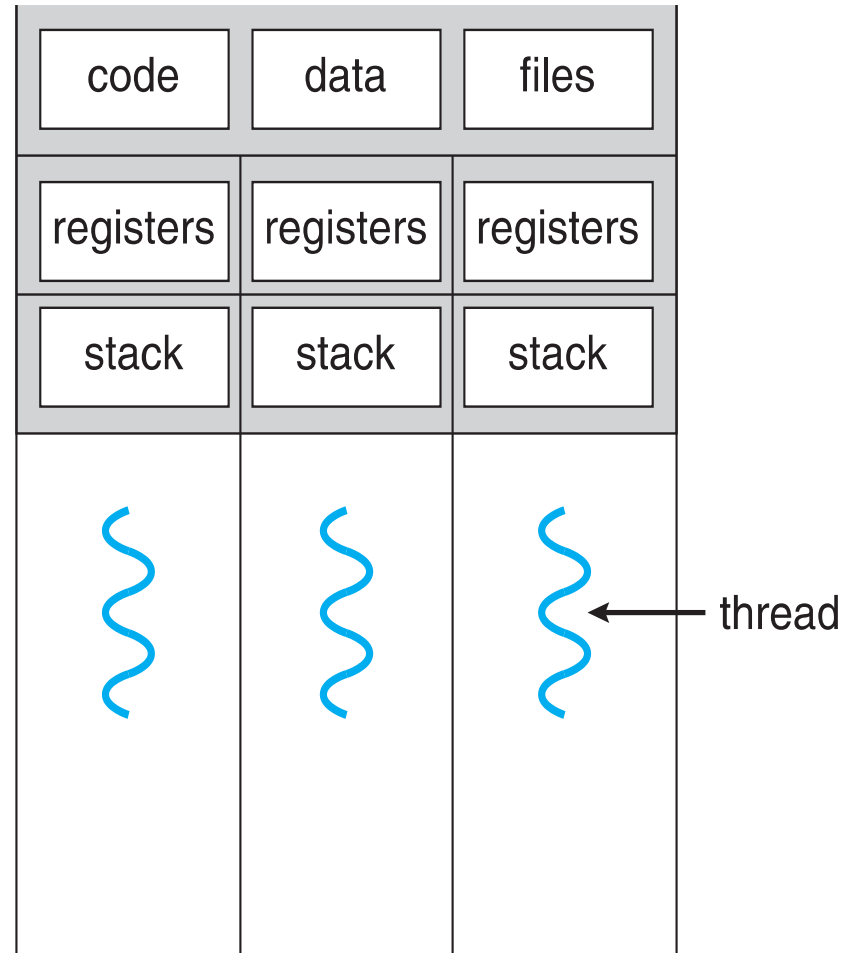
- Threads reside in the same address space of a single process
- All information exchange is via data shared between the threads
- Threads synchronize via simple primitives
- Each thread has its own stack, hardware registers, and state
- Thread table/switch: a lighter process table/switch
- Each thread may call any OS-supported system call on behalf of the process to which it belongs



Single and Multithreaded Processes



single-threaded process



multithreaded process

POSIX Threads

Thread call	Description
Pthread_create	Create a new thread
pthread_exit	Terminate the calling thread
pthread_join	Wait for a specific thread to exit
pthread_yield	Release the CPU to let another thread run
pthread_attr_init	Create and initialize a thread's attribute structure
pthread_attr_destroy	Remove a thread's attribute structure

Some of the P threads function calls.

Pthreads Example

What will the output be?

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

#define NUMBER_OF_THREADS 10
void * print_hello_world(void * tid)
{
    printf("Hello World. Greetings from thread %d\n", tid);
    pthread_exit(NULL);
}
int main(int argc, char * argv[])
{
    pthread_t threads[NUMBER_OF_THREADS];
    int status, i;
    for(i=0; i < NUMBER_OF_THREADS; i++) {
        status = pthread_create(&threads[i], NULL, print_hello_world, (void * )i);
        if (status != 0) {
            exit(-1);
        }
    }
    return 0;
}
```



Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```





User Threads and Kernel Threads


- ❑ **User threads** - management done by user-level threads library
- ❑ Three primary thread libraries:
 - ❑ POSIX **Pthreads**
 - ❑ Windows threads
 - ❑ Java threads
- ❑ **Kernel threads** - Supported by the Kernel
- ❑ Examples – virtually all general-purpose operating systems, including:
 - ❑ Windows
 - ❑ Solaris
 - ❑ Linux
 - ❑ Tru64 UNIX
 - ❑ Mac OS X

User and Kernel Level Threads

(a) A user-level threads package. (b) A threads package managed by the kernel.



User Threads: Pros and Cons



- + Thread switching time (no mode switch)
- + Scalability, customizability (no in-kernel management)
- Transparency (typically requires app cooperation)
- Parallelism (blocking syscalls are problematic)



Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many





Many-to-One


- ❑ Many user-level threads mapped to single kernel thread
- ❑ One thread blocking causes all to block
- ❑ Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- ❑ Few systems currently use this model
- ❑ Examples:
 - ❑ **Solaris Green Threads**
 - ❑ **GNU Portable Threads**





One-to-One


- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux
 - Solaris 9 and later





Many-to-Many Model


- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows with the *ThreadFiber* package





Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier





Implicit Threading

- ❑ Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- ❑ Creation and management of threads done by compilers and run-time libraries rather than programmers
- ❑ Three methods explored
 - ❑ Thread Pools
 - ❑ OpenMP
 - ❑ Grand Central Dispatch
- ❑ Other methods include Microsoft Threading Building Blocks (TBB), `java.util.concurrent` package





Thread Pools

- ❑ Create a number of threads in a pool where they await work
- ❑ Advantages:
 - ❑ Usually slightly faster to service a request with an existing thread than create a new thread
 - ❑ Allows the number of threads in the application(s) to be bound to the size of the pool
 - ❑ Separating task to be performed from mechanics of creating task allows different strategies for running task
 - ▶ i.e. Tasks could be scheduled to run periodically
- ❑ Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
     * this function runs as a separate thread.  
     */  
}
```





OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

#pragma omp parallel

Create as many threads as there are cores

```
#pragma omp parallel for
for(i=0;i<N;i++) {
    c[i] = a[i] + b[i];
}
```

Run for loop in parallel

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```





Grand Central Dispatch

- ❑ Apple technology for Mac OS X and iOS operating systems
- ❑ Extensions to C, C++ languages, API, and run-time library
- ❑ Allows identification of parallel sections
- ❑ Manages most of the details of threading
- ❑ Block is in “^{}” - `^ { printf("I am a block"); }`
- ❑ Blocks placed in dispatch queue
 - ❑ Assigned to available thread in thread pool when removed from queue



Threads: Issues

- Does the OS keep track of threads?
 - Kernel threads vs. user threads
- What to do on fork?
 - Clone all threads vs. calling thread
 - What if a thread is currently blocking on a systems call?
- What to do with signals?
 - Send signal to all threads vs. single thread
 - Per-process or per-thread signal handlers
- Where to store per-thread variables?
- Does sharing come at a cost?
- Are threads required inside an operating system?




Thread Cancellation

- ❑ Terminating a thread before it has finished
- ❑ Thread to be canceled is **target thread**
- ❑ Two general approaches:
 - ❑ **Asynchronous cancellation** terminates the target thread immediately
 - ❑ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- ❑ Pthread code to create and cancel a thread:

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
. . .  
  
/* cancel the thread */  
pthread_cancel(tid);
```



Fifth Edition



End of Threads Lecture